

Peer Response

Your post provides a strong and well-structured analysis of both the benefits and risks of AI writers, particularly your discussion on hallucinations and lack of true understanding in systems like GPT-3. I agree that these limitations pose serious risks in high-stakes domains such as healthcare and public policy.

To mitigate these issues, one important measure is the implementation of domain-restricted deployment. AI writing systems should be limited to specific use cases where the risks are lower (e.g., administrative drafting), while high-risk applications should require stricter validation protocols. This aligns with risk-based AI governance frameworks increasingly adopted in industry.

Another key solution is the integration of retrieval-augmented generation (RAG) approaches, where AI models are connected to verified external knowledge sources rather than relying solely on internal statistical patterns. This can significantly reduce hallucinations by grounding outputs in factual data, improving reliability in real-world applications.

Additionally, enhancing model evaluation and auditing processes is essential. Regular stress-testing of models using adversarial prompts can help identify weaknesses before deployment. From a Machine Learning perspective, this supports more robust and trustworthy systems.

Your point about bias is also critical. Beyond dataset improvements, organisations should adopt fairness monitoring tools to continuously detect and correct biased outputs after deployment (Khalifa & Albadawy, 2024).

Overall, your analysis effectively captures the key challenges, and combining technical improvements with governance strategies is essential to ensure AI writers are used safely and responsibly.



References (Harvard Style)

- Khalifa, M. and Albadawy, M. (2024) Using artificial intelligence in academic writing and research: An essential productivity tool. *Computer Methods and Programs in Biomedicine Update*, 5, 100145.