

Peer Response

Your discussion provides a clear and well-supported evaluation of the capabilities and risks of large language models such as GPT-3, particularly in highlighting their lack of true understanding despite impressive fluency. I strongly agree with your point regarding the dangers of deploying such systems in sensitive contexts, especially the example of harmful responses to mental health queries.

To address these risks, several measures could be implemented. Firstly, integrating human-in-the-loop oversight is critical, particularly in high-risk applications such as healthcare or legal advisory systems. AI-generated outputs should always be reviewed by qualified professionals before being delivered to users. This ensures accountability and reduces the likelihood of harmful or misleading information being disseminated.

Secondly, content filtering and alignment techniques should be strengthened during model development. Advances in Reinforcement Learning from Human Feedback have shown promise in reducing harmful outputs by aligning model behaviour with human values (Ouyang et al., 2022). However, continuous monitoring is still required, as no system is entirely risk-free.

Additionally, improving dataset curation and bias mitigation strategies is essential. As you noted, biased outputs often originate from training data, so implementing stricter data selection and auditing processes can help reduce this issue (Bender et al., 2021). Transparency measures, such as clearly labelling AI-generated content, can further enhance user trust and awareness.

Overall, your analysis effectively highlights the challenges of AI writers, and these combined technical and governance strategies are essential to ensure their safe and responsible use.



References (Harvard Style)

- Bender, E.M., Gebru, T., McMillan-Major, A. and Shmitchell, S. (2021) On the dangers of stochastic parrots: Can language models be too big? Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency.
- Ouyang, L. et al. (2022) Training language models to follow instructions with human feedback. arXiv preprint arXiv:2203.02155.