Peer Response to Abdulla Al-shibli

Abdullah, your synthesis highlights exactly where incidents arise in ACL-based MAS: semantic drift, brittle discovery, stalled negotiations, and latency under load. To prevent these, I'd institute strict ontology governance with machine-checkable constraints (e.g., SHACL) and versioned, deprecation-aware vocabularies so malformed or ambiguous messages are rejected early with actionable errors. Discovery should be resilient by replicating facilitator/directory agents, using health checks and signed capability manifests cached at the edge to survive outages. Interaction should be hardened by adopting standard FIPA protocols with explicit timeouts, retries with jitter, idempotency tokens, and caps on Contract Net rounds to avoid deadlocks and resource hoarding. For performance, use a hybrid policy: keep intent/coordination in ACL but route bulk/real-time data over RPC/streams, correlating both via conversation IDs for traceability. Add observability (structured logs by performative/ontology, distributed tracing, semantic anomaly alerts) and backpressure/QoS (rate limits, priority queues, graceful degradation). Secure by default with mutual auth, message signing, capability-scoped authorization, and envelope-level redaction. Finally, reduce change risk through contract tests for performatives, simulation of drops/reordering/duplication, and canary agents behind feature flags before full rollout. These measures preserve ACL expressiveness while systematically reducing the failure modes you noted (Finin et al., 1994; Labrou, Finin and Peng, 1999; FIPA, 2002; Bellifemine et al., 2008; W3C, 2017).

References

Bellifemine, F., Caire, G., Poggi, A. and Rimassa, G. (2008) 'JADE: A software framework for developing multi-agent applications', *Information and Software Technology*, 50(1–2), pp. 10–21. https://doi.org/10.1016/j.infsof.2007.10.008 (Accessed: 8 September 2025).
FIPA (2002) *FIPA Interaction Protocol Specifications*. Available at: https://www.fipa.org/repository/ (Accessed: 8 September 2025).
Finin, T., Fritzson, R., McKay, D. and McEntire, R. (1994) 'KQML as an agent communication language', *Proceedings of CIKM 'S4*, pp. 456–463. https://doi.org/10.1145/191246.191322 (Accessed: 8 September 2025).
Labrou, Y., Finin, T. and Peng, Y. (1999) 'Agent communication languages: The current landscape', *IEEE Intelligent Systems*, 14(2), pp. 45–52. https://doi.org/10.1109/5254.757631 (Accessed: 8 September 2025).
W3C (2017) *Shapes Constraint Language (SHACL)*. Available at: https://www.w3.org/TR/shacl/ (Accessed: 8 September 2025).

Peer Response to Fahad Abdallah

Fahad, to prevent the incidents you surface—semantic inconsistency, reasoning overload, and real-time performance hits—I'd harden MAS deployments along four fronts. First, enforce ontology governance: adopt a shared, versioned ontology with machine-checkable constraints (e.g., SHACL) and validation at ingress so malformed or ambiguous performatives are rejected early with actionable errors; include deprecation policies and mapping rules to curb drift (W3C, 2017). Second, constrain and standardise interactions: use a small, well-profiled subset of FIPA ACL performatives and documented FIPA interaction protocols with explicit timeouts, retries with jitter, and idempotency tokens; cap negotiation rounds (e.g., Contract Net) to avoid live lock and resource hoarding (FIPA, 2002; Labrou, Finin and Peng, 1999). Third, separate intent from throughput: keep ACL for tasking/negotiation but route bulk or latency-critical data over RPC/streams, correlating both via conversation IDs; add backpressure (rate limits, priority queues) and graceful degradation to keep real-time paths predictable. Fourth, make failure observable and safe: structured logs by performative/ontology, distributed tracing across agents, and semantic anomaly detection to flag contradictory or looping dialogues; pair with security controls (mutual auth, message signing, capability-scoped authorization) and replicated directories/facilitators with health checks and cached, signed capability manifests for resilience (Bellifemine et al., 2008). Finally, reduce reasoning cost and regressions with contract tests over message schemas, simulation harnesses that inject drops/reordering/duplication, and canary agents behind feature flags before full rollout. These measures preserve ACL flexibility while systematically addressing the real incidents you note—meaning fewer misunderstandings, fewer stalls, and predictable performance in open, heterogeneous environments (FIPA, 2002; Bellifemine et al., 2008; W3C, 2017; Labrou, Finin and Peng, 1999).

References

Bellifemine, F., Caire, G., Poggi, A. and Rimassa, G. (2008) 'JADE: A software framework for developing multi-agent applications. Lessons learned', *Information and Software Technology*, 50(1–2), pp. 10–21. Available at: https://doi.org/10.1016/j.infsof.2007.10.008 (Accessed: 8 September 2025).

FIPA (2002) *FIPA Interaction Protocol Specifications*. Available at: https://www.fipa.org/repository/ (Accessed: 8 September 2025).

Labrou, Y., Finin, T. and Peng, Y. (1999) 'Agent communication languages: The current landscape', *IEEE Intelligent Systems*, 14(2), pp. 45–52. Available at: https://doi.org/10.1109/5254.757631 (Accessed: 8 September 2025).

W3C (2017) *Shapes Constraint Language (SHACL)*. W3C Recommendation. Available at: https://www.w3.org/TR/shacl/ (Accessed: 8 September 2025).

Peer Response to Abdulla Alshaibani

Abdulla, to prevent the incidents you surface (semantic inconsistency, high reasoning overhead, brittle debugging), I'd institute a prevention stack that targets semantics, protocol discipline, resilience, and operability. First, ontology governance: use a shared, versioned ontology with machine-checkable constraints (e.g., SHACL) enforced at message ingress; add deprecation windows and mapping rules so agents can translate old terms, curbing drift and misunderstandings (Labrou, Finin and Peng, 1999; Afsharchi, Far and Denzinger, 2006). Second, lean ACL profiling: standardise on a small subset of performatives and documented interaction protocols, with explicit timeouts, retries with jitter, idempotency tokens, and hard caps on negotiation/Contract Net rounds to avoid livelock and resource hoarding (Gan Kim Soon et al., 2019; Hernández, Fallah-Seghrouchni and Soldano, 2004). Third, hybrid I/O lanes: keep ACL for intent and coordination but move bulk/latency-sensitive payloads to RPC/streams, correlating with conversation IDs; apply backpressure (rate limits, priority queues) and graceful degradation to preserve real-time predictability (Cooper and Torczon, 2012). Fourth, resilient discovery and trust: replicate facilitator/directory services with health checks; cache signed capability manifests so agents operate during outages. Finally, make failures observable before they bite: structured semantic logs (performatives, ontologies, conversation-ids), distributed tracing, dead-letter queues with quarantining, and simulation harnesses that inject drops/reordering/duplication; pair with contract tests over message schemas to catch regressions early. Together, these measures retain ACL flexibility where autonomy matters while preventing the common failure modes that push practitioners toward simpler method invocation in tightly coupled systems (Labrou, Finin and Peng, 1999; Cooper and Torczon, 2012).

References

Afsharchi, M., Far, B.H. and Denzinger, J. (2006) 'Ontology-guided learning to improve communication between groups of agents', *Proceedings of AAMAS 200c*, pp. 923–930. Available at: https://doi.org/10.1145/1160633.1160797 (Accessed: 8 September 2025).

Cooper, K.D. and Torczon, L. (2012) 'The procedure abstraction', in *Engineering a Compiler* (2nd ed.). Morgan Kaufmann, pp. 269–330. Available at: https://doi.org/10.1016/B978-0-12-088478-0.00006-2 (Accessed: 8 September 2025).

Gan, K.S., Chin, K.O., Anthony, P. and Hamdan, A.R. (2019) 'A review on agent communication language', in *Computational Science and Technology*. Springer, pp. 481–491. Available at: https://doi.org/10.1007/978-981-13-2622-6_47 (Accessed: 8 September 2025).

Hernández, A.G., Fallah-Seghrouchni, A.E. and Soldano, H. (2004) 'Distributed learning in intentional BDI multi-agent systems', *Proceedings of ENC 2004*, pp. 225–232. Available at: https://doi.org/10.1109/enc.2004.1342610 (Accessed: 8 September 2025).

Labrou, Y., Finin, T. and Peng, Y. (1999) 'Agent communication languages: The current landscape', *IEEE Intelligent Systems*, 14(2), pp. 45–52. Available at: https://doi.org/10.1109/5254.757631 (Accessed: 8 September 2025).

Peer Response to Rayyan Alnaqbi

Rayyan, to prevent the incidents you hint at—semantic mismatches, heavy reasoning costs, and slow paths—I'd put guardrails on four fronts. First, ontology governance and validation: agree a versioned shared ontology and enforce machine-checkable constraints (e.g., SHACL) at message ingress so malformed or out-of-date terms are rejected with actionable errors, reducing ambiguity early (Labrou, Finin and Peng, 1999; W3C, 2017). Second, protocol discipline: standardise on a minimal profile of FIPA performatives and interaction protocols with explicit timeouts, retries with jitter, idempotency tokens, and hard caps on negotiation rounds to avoid livelock and resource hoarding (FIPA, 2002; Wooldridge, 2009). Third, hybrid architecture: keep ACL for intent/coordination but route bulk or latency-critical data over RPC/streams, correlate both with conversation IDs, and apply backpressure (rate limits, priority queues) to keep real-time behaviour predictable (Wooldridge, 2009). Fourth, operability and resilience: replicate directory/facilitator services with health checks; cache signed capability manifests for offline operation; add structured semantic logs (performatives, ontology, conversation-id), distributed tracing, and dead-letter queues so failures are detectable and recoverable. Together, these measures retain ACLs' flexibility for autonomous cooperation while curbing the very overheads and inconsistencies that make teams default to simple method invocation in tightly coupled settings (Labrou, Finin and Peng, 1999; FIPA, 2002).

References

FIPA (2002) 'FIPA Communicative Act Library Specification'. Available at: https://www.fipa.org/specs/fipa00037/SC00037J.html (Accessed: 10 September 2025).

Labrou, Y., Finin, T. and Peng, Y. (1999) 'Agent communication languages: The current landscape', IEEE Intelligent Systems, 14(2), pp. 45–52. Available at: https://doi.org/10.1109/5254.757631 (Accessed: 10 September 2025).

W3C (2017) Shapes Constraint Language (SHACL). W3C Recommendation. Available at: https://www.w3.org/TR/shacl/ (Accessed: 10 September 2025).

Wooldridge, M. (2009) An Introduction to MultiAgent Systems. 2nd edn. Chichester: John Wiley & Sons. Available at: https://onlinelibrary.wiley.com/doi/book/10.1002/9781118720833 (Accessed: 10 September 2025).