

Linga, thank you for your clear and balanced discussion of ACLs versus method invocation. I found your point about ontologies particularly important: without a common vocabulary, even the most sophisticated communication protocol fails. This echoes a key challenge in multi-agent systems: semantic interoperability is as critical as the communication channel itself (Labrou, Finin & Peng, 1999).

Your emphasis on the overhead and verbosity of ACLs also resonated with my initial post. As you noted, the layered structure of KQML does make it less efficient than direct method invocation in Python or Java. In tightly coupled environments, the simplicity and speed of method calls remain unmatched. However, I would add that with advances in distributed middleware and lightweight frameworks, some of this overhead can now be mitigated — for example, through optimised parsing or modular ontologies that focus on domain-specific tasks.

One additional angle relates to feature representation, which we studied in Unit 7. The effectiveness of ACLs depends not only on the communication protocol but also on how knowledge is structured and exchanged. Even if the ontology is agreed upon, poorly engineered representations can reduce both interpretability and efficiency (Wooldridge, 2009). This highlights that the real challenge is not ACLs alone but the broader ecosystem of knowledge engineering.

Overall, I agree with your conclusion: ACLs are better suited to heterogeneous, distributed systems, while method invocation excels in monolithic or performance-driven contexts. Both approaches remind us that design choices should align with the intended domain, balancing efficiency, flexibility, and semantic richness.

References

- Labrou, Y., Finin, T. and Peng, Y. (1999) 'Agent communication languages: The current landscape', IEEE Intelligent Systems, 14(2), pp. 45–52.
- Wooldridge, M. (2009) An Introduction to MultiAgent Systems. 2nd edn. Chichester: Wiley.